



License Plate Detection and Modification

Group 7: Xiule Fan Shiva Tej Soma Xinyi Yao



Problem Formation



Practical Problem:


Recently, with the spread of car videos across the internet or television broadcasts, there are concerns about the privacy violation toward drivers since we can clearly see the vehicle license plate information through videos.

Importance of addressing above problem:


Addressing this problem can effectively protect vehicle license plate information and further protect the privacy of drivers.

Reason of Interest:

Our interest in solving this problem stems from the desire to promote privacy protection and facilitate video content creation without compromising the security and privacy of individuals.



Project Goal and Milestones




Project Goal:


We aim to develop a system that can detect and modify the license plates of cars in videos so that the license plate information will not be leaked.

Task breakdown:

- Milestone 1
 - License Plate Detection From Images
 - License Plate Detection From Videos
- Milestone 2
 - License Plate Modification In Images(inpainting/diffusion)



Datasets Used



YOLOv3:

- Kaggle dataset
- RoboFlow dataset
- 678 training images, 70 validation images, 35 test images

YOLOv5:

- Kaggle dataset of 473 images
- CCPD (Chinese City Parking Dataset)
 - This dataset consists of 250k+ images, including all kinds of augmentations.


CCPD (Chinese City Parking Dataset)

Edit


Introduced by Xu et al. in [Towards End-to-End License Plate Detection and Recognition: A Large Dataset and Baseline](#)

The **Chinese City Parking Dataset (CCPD)** is a dataset for license plate detection and recognition. It contains over 250k unique car images, with license plate location annotations.

Source:  [Towards End-to-End License Plate Detection and Recognition: A Large Dataset and Baseline](#)

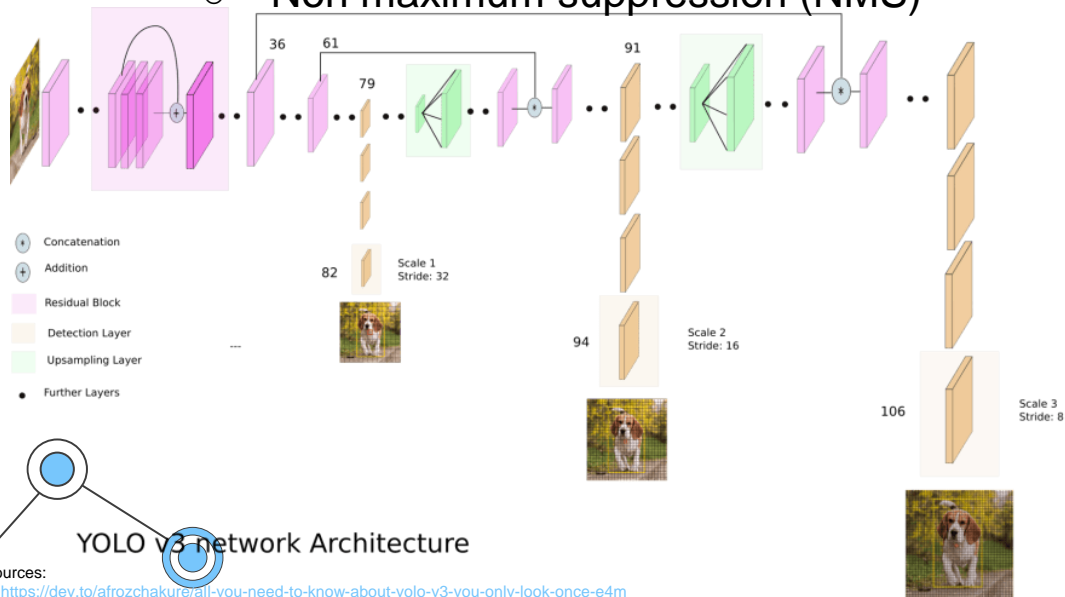


Source: Xu et al.




License Plate Detection from Images

- YOLOv3 - Model
 - Architecture
 - Non maximum suppression (NMS)




Sources:

- <https://dev.to/afrozchakure/all-you-need-to-know-about-yolo-v3-you-only-look-once-e4m>
- <https://www.analyticsvidhya.com/blog/2020/08/selecting-the-right-bounding-box-using-non-max-suppression-with-implementation/>




Multiple Bounding Boxes



Final Bounding Boxes

License Plate Detection from Images

- YOLOv3 - Result on images
 - Results: Average Precision (AP) = 0.7525




License Plate Detection from Images

- YOLOv5 - Result on images



On Kaggle Dataset

On CCPD2019 Dataset




License Plate Detection from Images

- YOLOv3 - Result on videos
 - False positive detections
 - Missing detections at some frames
- YOLOv5 - less consistent plate detections



License Plate Detection from Videos

- Seq-NMS
 - An extension of NMS for image sequences



Sources:

1. W. Han et al., "Seq-NMS for Video Object Detection," arXiv:1602.08465 [cs.CV], 2016.

License Plate Detection from Videos

- YOLOv3 with Seq-NMS
 - More consistent prediction throughout the sequence
 - False positive is still an issue




License Plate Detection and Modification

Click the
image to
open the
Video




MILESTONE 2



Blurring license plates may be an acceptable solution to protect privacy, but it can cause distraction and shift the focus from the main subject of the image. To address this issue, we conducted further research to find alternative solutions that would maintain the privacy of the individuals while also improving the overall aesthetics of the image.


We Tried Inpainting and Diffusion



Inpainting


Image Inpainting is a task of reconstructing missing regions in an image. It is an important problem in computer vision and an essential functionality in many imaging and graphics applications, e.g. object removal, image restoration, manipulation, re-targeting, compositing, and image-based rendering.

E.g: Healing tools in photoshop or
snapseed




Inpainting example, google photos

GAN Architectures for Inpainting




Used both CCPD and Kaggle datasets on below architecture variations (trained for 100 epochs each)

- Custom Generator and discriminator
- Pretrained U-Net with resnet encoder, custom discriminator
- Freezing few layers for pre-trained U-Net.



Custom GAN Architecture




Generator

Input Image (3 channels)
↓
Conv2d(3, 64) → ReLU
↓
Conv2d(64, 128) → BatchNorm2d → ReLU
↓
Conv2d(128, 256) → BatchNorm2d → ReLU
↓
ConvTranspose2d(256, 128) → BatchNorm2d → ReLU
↓
ConvTranspose2d(128, 64) → BatchNorm2d → ReLU
↓
ConvTranspose2d(64, 3) → Tanh
↓
Output Image (3 channels)

Discriminator

Input Image (3 channels)
↓
Conv2d(3, 64) → LeakyReLU(0.2)
↓
Conv2d(64, 128) → BatchNorm2d → LeakyReLU(0.2)
↓
Conv2d(128, 256) → BatchNorm2d → LeakyReLU(0.2)
↓
Conv2d(256, 1) → AdaptiveAvgPool2d → Sigmoid
↓
Output (1 channel)



Generator

Input Image (3 channels)

↓

Conv2d(3, 64) → ReLU

↓

Conv2d(64, 128) → BatchNorm2d → ReLU

↓

Conv2d(128, 256) → BatchNorm2d → ReLU

↓

ConvTranspose2d(256, 128) → BatchNorm2d → ReLU

↓

ConvTranspose2d(128, 64) → BatchNorm2d → ReLU

↓

ConvTranspose2d(64, 3) → Tanh

↓

Output Image (3 channels)

Discriminator

Input Image (3 channels)

↓

Conv2d(3, 64) → LeakyReLU(0.2)

↓

Conv2d(64, 128) → BatchNorm2d → LeakyReLU(0.2)

↓

Conv2d(128, 256) → BatchNorm2d → LeakyReLU(0.2)


↓

Conv2d(256, 1) → AdaptiveAvgPool2d → Sigmoid


↓

Output (1 channel)

Observations



- All models resulted in Blurry outputs
- According to the Research paper from 2016 (University of California) L2 loss prefers blurred images therefore it is suggested to divide losses into reconstruction loss and adversarial loss




Context Encoders: Feature Learning by Inpainting

Deepak Pathak Philipp Krähenbühl Jeff Donahue Trevor Darrell Alexei A. Efros
University of California, Berkeley
{pathak, philkr, jdonahue, trevor, efros}@cs.berkeley.edu

Abstract

We present an unsupervised visual feature learning algorithm driven by context-based pixel prediction. By analogy with auto-encoders, we propose Context Encoders – a convolutional neural network trained to generate the contents of an arbitrary image region conditioned on its surroundings. In order to succeed at this task, context encoders need to both understand the content of the entire image, as well as produce a plausible hypothesis for the missing part(s). When training context encoders, we have experimented with both a standard pixel-wise reconstruction loss, as well as a reconstruction plus an adversarial loss. The latter produces much sharper results because it can better handle multiple modes in the output. We found that a context encoder learns a representation that captures not just appearance but also the semantics of visual structures. We quantitatively demonstrate the effectiveness of our learned features for CNN pre-training on classification, detection, and segmentation tasks. Furthermore, context encoders can be used for semantic inpainting tasks, either stand-alone or as initialization for non-parametric methods.

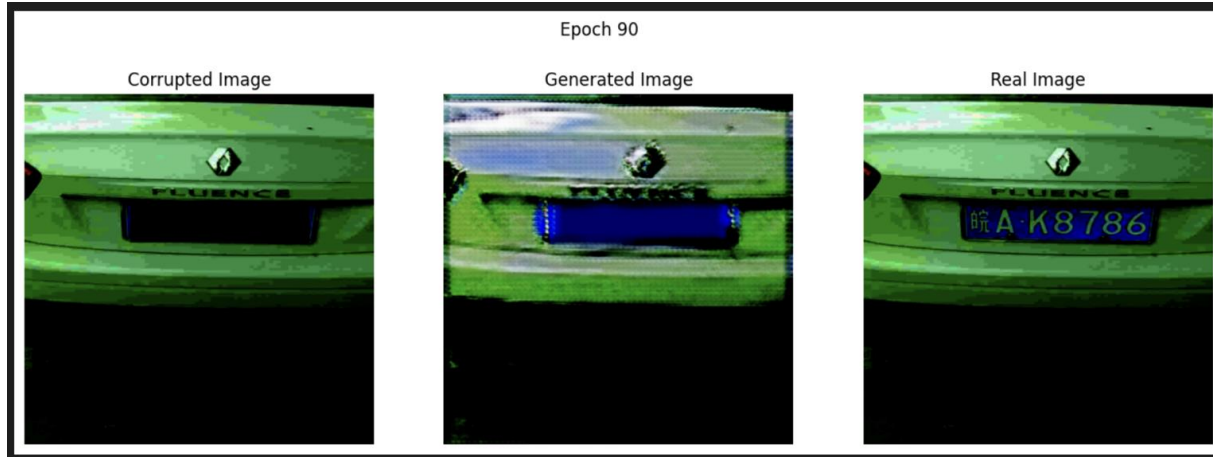


14.07379v2 [cs.CV] 21 Nov 2016

[1604.07379.pdf \(arxiv.org\)](https://arxiv.org/pdf/1604.07379.pdf)

Observations

- All of the Algorithms did an ok job inpainting, but at the cost of overall decrease in quality
- Most often the they are just filled with dominant color.



Custom GAN output

Observations

- When used pre-Trained U-Net with ResNet34 as encoder

Which one is generated?



Observations

- When used pre-Trained U-Net with ResNet34 as encoder

Real



Generated




Generated image colors looks realistic

Diffusion


Diffusion:

- Removes data by slowly adding gaussian blur
- Train a neural network to reverse the corruption process.
- It's a fascinating approach that has proved quite effective in Image Generation tasks
- E.g: DALLÉ2, Midjourney, Stable Diffusion



Diffusion inpainting example


Idea



Real Image




Idea




Real Image




Corrupted Image




Idea



Real Image



Corrupted Image

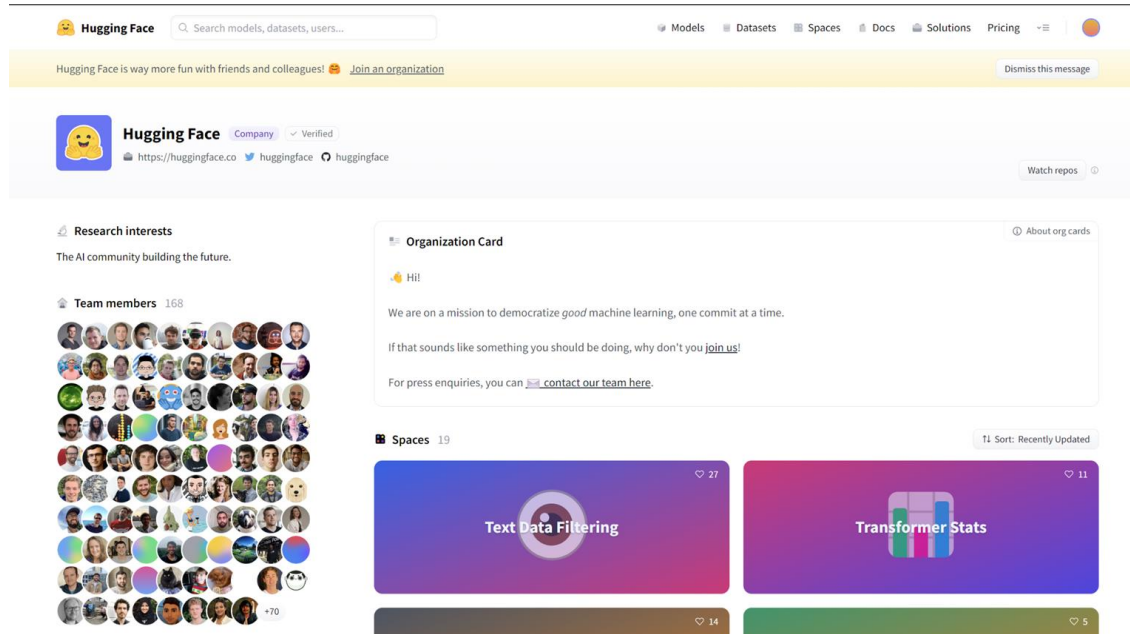


Generated Image



Tool used: Hugging face (Diffuser)

Consists of open source AI models and Tools to build them | diffuser module (2022)



The screenshot displays the Hugging Face website. At the top, there is a navigation bar with a search bar and links for Models, Datasets, Spaces, Docs, Solutions, and Pricing. Below this is a yellow banner with the text "Hugging Face is way more fun with friends and colleagues!" and a link to "Join an organization". The main content area features the Hugging Face logo, a verified company badge, and links to their website and social media. A "Research interests" section describes the community's goal. A "Team members" section shows a grid of 168 member avatars. An "Organization Card" provides a mission statement and contact information. A "Spaces" section lists featured content, including "Text Data Filtering" and "Transformer Stats".

Hugging Face Search models, datasets, users... Models Datasets Spaces Docs Solutions Pricing

Hugging Face is way more fun with friends and colleagues! [Join an organization](#) Dismiss this message

Hugging Face Company Verified <https://huggingface.co> [huggingface](#) [huggingface](#) Watch repos

Research interests
The AI community building the future.

Team members 168

Organization Card About org cards
Hi!
We are on a mission to democratize *good* machine learning, one commit at a time.
If that sounds like something you should be doing, why don't you [join us!](#)
For press enquiries, you can [contact our team here](#).

Spaces 19 T1 Sort: Recently Updated

Text Data Filtering 27


Transformer Stats 11

Results trained on license plates alone



Text is not legible

Discussion




What we have learned from this project:

- Data augmentation
- Thresholds for different parameters
- Anchor box size
- Consumer Hardware Limitations
- Reading Research Paper
- Trick: Leverage parallel computing for cpu bound tasks (data preparation/augmentation)



PC crash when training



Thanks!

Many more test videos and Notebook file in form of tutorial are available.
Feel free to ask



License Plate Detection from Videos

- YOLOv3 with Seq-NMS

